

```
In [92]: #Load the librarys
import pandas as pd #To work with dataset
import numpy as np #Math library
import seaborn as sns #Graph library that use matplotlib in background
import matplotlib.pyplot as plt #to plot some parameters in seaborn
import warnings

warnings.filterwarnings('ignore')

#Importing the data
df_credit = pd.read_csv("/Users/wandawu/Desktop/批借贷问题/german_credit_data")
```

```
In [93]: df_credit.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1000 entries, 0 to 999
Data columns (total 11 columns):
Unnamed: 0    1000 non-null int64
Age           1000 non-null int64
Sex           1000 non-null object
Job           1000 non-null int64
Housing       1000 non-null object
Saving accounts 817 non-null object
Checking account 606 non-null object
Credit amount 1000 non-null int64
Duration      1000 non-null int64
Purpose       1000 non-null object
Risk          1000 non-null object
dtypes: int64(5), object(6)
memory usage: 86.0+ KB
```

```
In [94]: df_credit.head()
```

Out[94]:

	Unnamed: 0	Age	Sex	Job	Housing	Saving accounts	Checking account	Credit amount	Duration	Purpose
0	0	67	male	2	own	NaN	little	1169	6	radio/T
1	1	22	female	2	own	little	moderate	5951	48	radio/T
2	2	49	male	1	own	little	NaN	2096	12	educatio
3	3	45	male	2	free	little	little	7882	42	furniture/equipmer
4	4	53	male	2	free	little	little	4870	24	ca


```
In [95]: df_credit.nunique()
```

```
Out[95]: Unnamed: 0      1000  
Age      53  
Sex      2  
Job      4  
Housing  3  
Saving accounts  4  
Checking account  3  
Credit amount  921  
Duration  33  
Purpose  8  
Risk     2  
dtype: int64
```

```
In [96]: interval = (18, 25, 35, 60, 120)  
  
cats = ['Student', 'Young', 'Adult', 'Senior']  
df_credit["Age_cat"] = pd.cut(df_credit.Age, interval, labels=cats)
```


```
In [97]: plt.subplot2grid((1,2),(0,0))  
df_credit.Risk.value_counts().plot(kind='bar')  
plt.ylabel(u'number')  
  
plt.subplot2grid((1,2),(0,1))  
df_credit.Age.plot(kind='kde')
```

```
Out[97]: <matplotlib.axes._subplots.AxesSubplot at 0x1a1f50a080>
```



```
In [98]: plt.subplot(121)
df_credit.Age[df_credit.Risk == 'good'].plot(kind='kde')
plt.title(u'good')
plt.subplot(122)
df_credit.Age[df_credit.Risk == 'bad'].plot(kind='kde')
plt.title(u'bad')
```

Out[98]: Text(0.5,1,'bad')



```
In [99]: interval = (18, 25, 35, 60, 120)

label = ['Student', 'Adult', 'Senior', 'Old']
df_credit_age = pd.cut(df_credit.Age, interval, labels=label)

df_good = df_credit.Age_cat[df_credit.Risk == 'good'].value_counts()
df_bad = df_credit.Age_cat[df_credit.Risk == 'bad'].value_counts()
df_credit.head()
```


Out[99]:

	Unnamed: 0	Age	Sex	Job	Housing	Saving accounts	Checking account	Credit amount	Duration	Purpose
0	0	67	male	2	own	NaN	little	1169	6	radio/TV
1	1	22	female	2	own	little	moderate	5951	48	radio/TV
2	2	49	male	1	own	little	NaN	2096	12	education
3	3	45	male	2	free	little	little	7882	42	furniture/equipment
4	4	53	male	2	free	little	little	4870	24	car

```
In [100]: #look the credit risk in different age interval

df=pd.DataFrame({u'good':df_good, u'bad':df_bad})
df.plot(kind='bar',stacked=True)
#looks like the higher age the lower percentage of bad credit, it make sense
```

```
Out[100]: <matplotlib.axes._subplots.AxesSubplot at 0x1a1f56fb38>
```



```
In [101]: #now look at the sex
fig = plt.figure()
fig.set(alpha=0.2)


Sex_g = df_credit.Sex[df_credit.Risk == 'good'].value_counts()
Sex_b = df_credit.Sex[df_credit.Risk == 'bad'].value_counts()

df=pd.DataFrame({u'good':Sex_g,u'bad':Sex_b})
df.plot(kind='bar',stacked=True)

df.to_csv('test.csv',index=True,header=True,sep=",")

#seems like that the female has higher bad risk ratio?
```

<Figure size 432x288 with 0 Axes>



```
In [102]: # get dummies of categorical variable
df_credit.head()
```

Out[102]:

	Unnamed: 0	Age	Sex	Job	Housing	Saving accounts	Checking account	Credit amount	Duration	Purpose
0	0	67	male	2	own	NaN	little	1169	6	radio/T
1	1	22	female	2	own	little	moderate	5951	48	radio/T
2	2	49	male	1	own	little	NaN	2096	12	educatio
3	3	45	male	2	free	little	little	7882	42	furniture/equipmer
4	4	53	male	2	free	little	little	4870	24	ca

```
In [103]: dummies_sex = pd.get_dummies(df_credit.Sex,prefix='Sex')
dummies_Housing = pd.get_dummies(df_credit.Housing,prefix='Housing')
dummies_Saving = pd.get_dummies(df_credit["Saving accounts"],prefix='Saving')
dummies_Checking = pd.get_dummies(df_credit["Checking account"],prefix='Checking')
dummies_Purpose = pd.get_dummies(df_credit.Purpose,prefix='Purpose')
dummies_Risk = pd.get_dummies(df_credit.Risk,prefix='Risk')
df_1 = pd.concat([df_credit,dummies_sex,dummies_Saving,dummies_Housing,dummies_Purpose,dummies_Risk])
```

```
In [104]: df_1.head()
```

Out[104]:

	Unnamed: 0	Age	Sex	Job	Housing	Saving accounts	Checking account	Credit amount	Duration	Purpose
0	0	67	male	2	own	NaN	little	1169	6	radio/TV
1	1	22	female	2	own	little	moderate	5951	48	radio/TV
2	2	49	male	1	own	little	NaN	2096	12	education
3	3	45	male	2	free	little	little	7882	42	furniture/equipment
4	4	53	male	2	free	little	little	4870	24	car

5 rows × 34 columns

```
In [105]: df_1.drop(['Sex','Housing','Saving accounts','Checking account','Purpose','Risk'])
```

```
In [106]: df_1.head()
```

Out[106]:

	Unnamed: 0	Age	Job	Credit amount	Duration	Sex_female	Sex_male	Saving_little	Saving_moderate
0	0	67	2	1169	6	0	1	0	0
1	1	22	2	5951	48	1	0	1	0
2	2	49	1	2096	12	0	1	1	0
3	3	45	2	7882	42	0	1	1	0
4	4	53	2	4870	24	0	1	1	0

5 rows × 26 columns

```
In [107]: from sklearn.model_selection import train_test_split, KFold, cross_val_score
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report

from sklearn.model_selection import GridSearchCV

# Algorithms models to be compared
from sklearn.ensemble import RandomForestClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.discriminant_analysis import LinearDiscriminantAnalysis
from sklearn.naive_bayes import GaussianNB
from sklearn.svm import SVC
from xgboost import XGBClassifier
```

```
In [108]: # scaling age, credit amount, duration
import sklearn.preprocessing as preprocessing
scaler = preprocessing.StandardScaler()
df_1['Age_scaled'] = scaler.fit_transform(df_1['Age'].values.reshape(-1,1))
df_1['Creditamount_scaled'] = scaler.fit_transform(df_1['Credit amount'].values.reshape(-1,1))
df_1['Duration_scaled'] = scaler.fit_transform(df_1['Duration'].values.reshape(-1,1))
```

```
In [109]: #select features we need from df_1

df_2 = df_1.filter(regex='Risk_good|Age_scaled|Job|Creditamount_scaled|Duration_scaled')
```

```
In [110]: x = df_2.drop('Risk_good', 1).values
y = df_2['Risk_good'].values
```

```
In [111]: x_train, x_test, y_train, y_test = train_test_split(x, y, test_size= 0.25, random_state= 42)
```

```

In [112]: # prepare models
models = []
models.append(('LR', LogisticRegression()))
models.append(('LDA', LinearDiscriminantAnalysis()))
models.append(('KNN', KNeighborsClassifier()))
models.append(('CART', DecisionTreeClassifier()))
models.append(('NB', GaussianNB()))
models.append(('RF', RandomForestClassifier()))
models.append(('SVM', SVC(gamma='auto')))
models.append(('XGB', XGBClassifier()))

# evaluate each model in turn
results = []
names = []
scoring = 'recall'

for name, model in models:
    kfold = KFold(n_splits=10, random_state=1)
    cv_results = cross_val_score(model, x_train, y_train, cv=kfold, scoring=scoring)
    results.append(cv_results)
    names.append(name)
    msg = "%s: %f (%f)" % (name, cv_results.mean(), cv_results.std())
    print(msg)

# boxplot algorithm comparison
fig = plt.figure(figsize=(11,6))
fig.suptitle('Algorithm Comparison')
ax = fig.add_subplot(111)
plt.boxplot(results)
ax.set_xticklabels(names)
plt.show()

```


```

LR: 0.885711 (0.037889)
LDA: 0.853954 (0.043644)
KNN: 0.859425 (0.043830)
CART: 0.743592 (0.055031)
NB: 0.732933 (0.045530)
RF: 0.806878 (0.073234)
SVM: 0.977493 (0.022456)
XGB: 0.885911 (0.032338)

```



## Algorithm Comparison



```
In [113]: # XGB model
# test set accuracy
XGB = XGBClassifier()

# Fitting with train data
model_XGB = XGB.fit(x_train, y_train)
# Printing the Training Score
print("Training score data: ")
print(model_XGB.score(x_train, y_train))
```

Training score data:  
0.8413333333333334

```
In [114]: y_pred = model_XGB.predict(x_test)

print(accuracy_score(y_test,y_pred))
print("\n")
print(confusion_matrix(y_test, y_pred))
print("\n")
print(classification_report(y_test, y_pred))
```

0.732

```
[[ 26  48]
 [ 19 157]]
```

	precision	recall	f1-score	support
0	0.58	0.35	0.44	74
1	0.77	0.89	0.82	176
micro avg	0.73	0.73	0.73	250
macro avg	0.67	0.62	0.63	250
weighted avg	0.71	0.73	0.71	250

```
In [24]: # try random forest?
#Setting the Hyper Parameters
param_grid = {"max_depth": [3,5, 7, 10, None],
              "n_estimators": [3,5,10,25,50,150,500],
              "max_features": [4,7,15,20]}

#Creating the classifier
model_RF = RandomForestClassifier(random_state=2)

grid_search = GridSearchCV(model_RF, param_grid=param_grid, cv=5, scoring='r
grid_search.fit(x_train, y_train)
print(grid_search.best_score_)
print(grid_search.best_params_)
```

Fitting 5 folds for each of 140 candidates, totalling 700 fits

```
[CV] max_depth=3, max_features=4, n_estimators=3 .....
[CV]  max_depth=3, max_features=4, n_estimators=3, score=1.0, total=
0.0s
[CV] max_depth=3, max_features=4, n_estimators=3 .....
[CV]  max_depth=3, max_features=4, n_estimators=3, score=0.990476190476
1905, total=    0.0s
[CV] max_depth=3, max_features=4, n_estimators=3 .....
[CV]  max_depth=3, max_features=4, n_estimators=3, score=0.895238095238
0953, total=    0.0s
[CV] max_depth=3, max_features=4, n_estimators=3 .....
[CV]  max_depth=3, max_features=4, n_estimators=3, score=0.942857142857
1428, total=    0.0s
[CV] max_depth=3, max_features=4, n_estimators=3 .....
[CV]  max_depth=3, max_features=4, n_estimators=3, score=0.961538461538
4616, total=    0.0s
[CV] max_depth=3, max_features=4, n_estimators=5 .....
[CV]  max_depth=3, max_features=4, n_estimators=5, score=1.0, total=
0.0s
-----
```

```
In [25]: rf = RandomForestClassifier(max_depth=3, max_features=4, n_estimators=150,
#training with the best params
rf.fit(x_train, y_train)
```

```
Out[25]: RandomForestClassifier(bootstrap=True, class_weight=None, criterion='gin
i',
                                max_depth=3, max_features=4, max_leaf_nodes=None,
                                min_impurity_decrease=0.0, min_impurity_split=None,
                                min_samples_leaf=1, min_samples_split=2,
                                min_weight_fraction_leaf=0.0, n_estimators=150, n_jobs=None,
                                oob_score=False, random_state=2, verbose=0, warm_start=False)
```

```
In [115]: #Testing the model
#Predicting using our model
y_pred = rf.predict(x_test)

# Verificaar os resultados obtidos
print(accuracy_score(y_test,y_pred))
print("\n")
print(confusion_matrix(y_test, y_pred))
print("\n")
print(fbeta_score(y_test, y_pred, beta=2))
```

```
-----
--
ValueError                                Traceback (most recent call last)
<ipython-input-115-ed95976e639d> in <module>()
      1 #Testing the model
      2 #Predicting using our model
----> 3 y_pred = rf.predict(x_test)
      4
      5 # Verificaar os resultados obtidos

<ipython-input-74-5eacca0da31e> in predict(self, x)
      9
     10     def predict(self, x):
----> 11         return self.clf.predict(x)
     12
     13     def fit(self,x,y):

/anaconda3/lib/python3.6/site-packages/sklearn/ensemble/forest.py in predict
(self, X)
     541         The predicted classes.
     542         """
--> 543         proba = self.predict_proba(X)
     544
     545         if self.n_outputs_ == 1:

/anaconda3/lib/python3.6/site-packages/sklearn/ensemble/forest.py in predict
_proba(self, X)
     581         check_is_fitted(self, 'estimators_')
     582         # Check data
--> 583         X = self._validate_X_predict(X)
     584
     585         # Assign chunk of trees to jobs

/anaconda3/lib/python3.6/site-packages/sklearn/ensemble/forest.py in _val
idate_X_predict(self, X)
     360                                     "call `fit` before exploiting the
model.")
     361
--> 362         return self.estimators_[0]._validate_X_predict(X, check_input=True)
     363
     364     @property

/anaconda3/lib/python3.6/site-packages/sklearn/tree/tree.py in _validate_X
_predict(self, X, check_input)
```

```

386                                     "match the input. Model n_features i
s %s and "
387                                     "input n_features is %s "
--> 388                                     % (self.n_features_, n_features))
389
390                                     return X

```

**ValueError:** Number of features of the model must match the input. Model n\_features is 5 and input n\_features is 24

```

In [116]: #SVM
# test set accuracy
SVM = SVC(gamma='auto')

# Fitting with train data
model_SVM = SVM.fit(x_train, y_train)
# Printing the Training Score
print("Training score data: ")
print(model_SVM.score(x_train, y_train))

```

Training score data:  
0.752

```

In [117]: y_pred = model_SVM.predict(x_test)

print(accuracy_score(y_test,y_pred))
print("\n")
print(confusion_matrix(y_test, y_pred))
print("\n")
print(classification_report(y_test, y_pred))

```

0.724

```

[[ 11  63]
 [  6 170]]

```

	precision	recall	f1-score	support
0	0.65	0.15	0.24	74
1	0.73	0.97	0.83	176
micro avg	0.72	0.72	0.72	250
macro avg	0.69	0.56	0.54	250
weighted avg	0.71	0.72	0.66	250

```
In [118]: # GaussianNB
GNB = GaussianNB()

# Fitting with train data
model = GNB.fit(x_train, y_train)
# Printing the Training Score
print("Training score data: ")
print(model.score(x_train, y_train))
```

Training score data:  
0.6826666666666666

```
In [119]: y_pred = model.predict(x_test)

print(accuracy_score(y_test, y_pred))
print("\n")
print(confusion_matrix(y_test, y_pred))
print("\n")
print(classification_report(y_test, y_pred))
```

0.632

```
[[ 39  35]
 [ 57 119]]
```

	precision	recall	f1-score	support
0	0.41	0.53	0.46	74
1	0.77	0.68	0.72	176
micro avg	0.63	0.63	0.63	250
macro avg	0.59	0.60	0.59	250
weighted avg	0.66	0.63	0.64	250

```
In [120]: #xgboost XGBClassifier has best performance
```

```
In [121]: # let's do model stacking!
train= np.column_stack((x_train,y_train))
test = np.column_stack((x_test, y_test))
```

```
In [122]: LR= LogisticRegression()
LDA= LinearDiscriminantAnalysis()
KNN = KNeighborsClassifier()
CART = DecisionTreeClassifier()
NB = GaussianNB()
SVM = SVC(gamma='auto')
XGB = XGBClassifier()
RF = RandomForestClassifier()
```

In [34]:

```
(750, 24) (750,) (250, 24)
```

In [124]:

```
x = df_2.drop('Risk_good', 1)
y = df_2['Risk_good']
x_1, x_2, y_1, y_2 = train_test_split(x, y, test_size= 0.25, random_state=1)
train = x_1
ntrain = train.shape[0]  ## train set number
ntest = test.shape[0]    ## test set number
print(x_train.shape, y_train.shape, x_test.shape)
```

```
from sklearn.model_selection import KFold,StratifiedKFold
SEED = 0 # for reproducibility

NFOLDS = 5 # set folds for out-of-fold prediction
kf = KFold(n_splits= NFOLDS, shuffle=False,random_state=SEED)
```

```
(750, 24) (750,) (250, 24)
```

In [74]:

```
# Class to extend the Sklearn classifier
class SklearnHelper(object):
    def __init__(self, clf, seed=0, params=None):
        params['random_state'] = seed
        self.clf = clf(**params)

    def train(self, x_train, y_train):
        self.clf.fit(x_train, y_train)

    def predict(self, x):
        return self.clf.predict(x)

    def fit(self,x,y):
        return self.clf.fit(x,y)

    def feature_importances(self,x,y):
        print(self.clf.fit(x,y).feature_importances_)

# Class to extend XGboost classifier
```

```
In [75]: #Out-of-Fold Predictions
def get_oof(clf, x_train, y_train, x_test):
    oof_train = np.zeros((ntrain,))
    oof_test = np.zeros((ntest,))
    oof_test_skf = np.empty((NFOLDS, ntest))

    for i, (train_index, test_index) in enumerate(kf.split(train)):
        x_tr = x_train[train_index]
        y_tr = y_train[train_index]
        x_te = x_train[test_index]

        clf.train(x_tr, y_tr)

        oof_train[test_index] = clf.predict(x_te)
        oof_test_skf[i, :] = clf.predict(x_test)

    oof_test[:] = oof_test_skf.mean(axis=0)
    return oof_train.reshape(-1, 1), oof_test.reshape(-1, 1)
```



```
In [76]: #Generating our Base First-Level Models¶
# Put in our parameters for said classifiers
# Random Forest parameters
rf_params = {
    'n_jobs': -1,
    'n_estimators': 500,
    'warm_start': True,
    #'max_features': 0.2,
    'max_depth': 3,
    'min_samples_leaf': 2,
    'max_features' : 'sqrt',
    'verbose': 0
}

# KNN
et_params = {
    'n_neighbors': 15 ,
    'weights': 'distance'
}

# xgb
ada_params = {
    'n_estimators': 500,
    'max_depth': 4,
    'min_child_weight': 2,
    #gamma=1,
    'gamma':0.9,
    'subsample':0.8,
    'colsample_bytree':0.8,
    'objective': 'binary:logistic',
    'nthread': -1,
    'scale_pos_weight':1
}

# LR
gb_params = {
    'C':1e5
}

# Support Vector Classifier parameters
svc_params = {
    'kernel' : 'linear',
    'C' : 0.025
}
```

In [77]:

```
# Create 5 objects that represent our 4 models
rf = SklearnHelper(clf=RandomForestClassifier, seed=SEED, params=rf_params)
et = SklearnHelper(clf=KNeighborsClassifier, seed=SEED, params=et_params)
ada = SklearnHelper(clf=XGBClassifier, seed=SEED, params=ada_params)
gb = SklearnHelper(clf=LogisticRegression, seed=SEED, params=gb_params)
svc = SklearnHelper(clf=SVC, seed=SEED, params=svc_params)
```

```
-----
--
TypeError                                Traceback (most recent call last)
<ipython-input-77-b720851ebaed> in <module>()
      2 # Create 5 objects that represent our 4 models
      3 rf = SklearnHelper(clf=RandomForestClassifier, seed=SEED, params=
rf_params)
----> 4 et = SklearnHelper(clf=KNeighborsClassifier, seed=SEED, params=et
_params)
      5 ada = SklearnHelper(clf=XGBClassifier, seed=SEED, params=ada_para
ms)
      6 gb = SklearnHelper(clf=LogisticRegression, seed=SEED, params=gb_p
arams)

<ipython-input-74-5eeca0da31e> in __init__(self, clf, seed, params)
      3     def __init__(self, clf, seed=0, params=None):
      4         params['random_state'] = seed
----> 5         self.clf = clf(**params)
      6
      7     def train(self, x_train, y_train):

/anaconda3/lib/python3.6/site-packages/sklearn/neighbors/classification.p
y in __init__(self, n_neighbors, weights, algorithm, leaf_size, p, metri
c, metric_params, n_jobs, **kwargs)
    128         leaf_size=leaf_size, metric=metric, p=p,
    129         metric_params=metric_params,
--> 130         n_jobs=n_jobs, **kwargs)
    131         self.weights = _check_weights(weights)
    132

TypeError: __init__() got an unexpected keyword argument 'random_state'
```

In [79]:

```
et_oof_train, et_oof_test = get_oof(et, x_train, y_train, x_test) # Extra T
rf_oof_train, rf_oof_test = get_oof(rf, x_train, y_train, x_test) # Random F
ada_oof_train, ada_oof_test = get_oof(ada, x_train, y_train, x_test) # AdaB
gb_oof_train, gb_oof_test = get_oof(gb, x_train, y_train, x_test) # Gradient
svc_oof_train, svc_oof_test = get_oof(svc, x_train, y_train, x_test) # Suppor
print("Training is complete")
```

Training is complete

```
In [80]: rf_feature = rf.feature_importances(x_train,y_train)
et_feature = et.feature_importances(x_train, y_train)
ada_feature = ada.feature_importances(x_train, y_train)
gb_feature = gb.feature_importances(x_train,y_train)
```

```
[0.03148122 0.26751512 0.28554869 0.14185021 0.27360476]
[0.00284214 0.31221472 0.27907916 0.07390442 0.33195956]
[0.48666667 0.16          0.14666667 0.09333333 0.11333333]
[0.00285865 0.24994616 0.70857532 0.00324428 0.03537558]
```

```
In [81]: rf_features = [0.03148122, 0.26751512, 0.28554869, 0.14185021, 0.27360476]
et_features = [0.00284214, 0.31221472, 0.27907916, 0.07390442, 0.33195956]
ada_features = [0.48666667, 0.16,          0.14666667, 0.09333333, 0.11333333]
gb_features = [0.00285865, 0.24994616, 0.70857532, 0.00324428, 0.03537558]
```

```
In [82]: x = df_2.drop('Risk_good', 1)
y = df_2['Risk_good']
x_1, x_2, y_1, y_2 = train_test_split(x, y, test_size= 0.25, random_state=1)
train = x_1
```

```
In [84]: cols = train.columns.values
# Create a dataframe with features
feature_dataframe = pd.DataFrame( {'features': cols,
    'Random Forest feature importances': rf_features,
    'Extra Trees feature importances': et_features,
    'AdaBoost feature importances': ada_features,
    'Gradient Boost feature importances': gb_features
})
```

```
-----
--
ValueError                                Traceback (most recent call last)
<ipython-input-84-9370a797e020> in <module>()
      5     'Extra Trees feature importances': et_features,
      6     'AdaBoost feature importances': ada_features,
----> 7     'Gradient Boost feature importances': gb_features
      8 })

/anaconda3/lib/python3.6/site-packages/pandas/core/frame.py in __init__(self, data, index, columns, dtype, copy)
    346                                     dtype=dtype, copy=copy)
    347     elif isinstance(data, dict):
--> 348         mgr = self._init_dict(data, index, columns, dtype=dtype)
    349     elif isinstance(data, ma.MaskedArray):
    350         import numpy.ma.mrecords as mrecords

/anaconda3/lib/python3.6/site-packages/pandas/core/frame.py in _init_dict(self, data, index, columns, dtype)
    457         arrays = [data[k] for k in keys]
    458
--> 459         return _arrays_to_mgr(arrays, data_names, index, columns,
    460                               dtype=dtype)
    461     def _init_ndarray(self, values, index, columns, dtype=None, copy=False):

/anaconda3/lib/python3.6/site-packages/pandas/core/frame.py in _arrays_to_mgr(arrays, arr_names, index, columns, dtype)
    7313     # figure out the index, if necessary
    7314     if index is None:
-> 7315         index = extract_index(arrays)
    7316
    7317     # don't force copy because getting jammed in an ndarray anyway

/anaconda3/lib/python3.6/site-packages/pandas/core/frame.py in extract_index(data)
    7359         lengths = list(set(raw_lengths))
    7360         if len(lengths) > 1:
-> 7361             raise ValueError('arrays must all be same length')
    7362
    7363         if have_dicts:
```



```
In [89]: y_pred = predictions
print(accuracy_score(y_test,y_pred))
print("\n")
print(confusion_matrix(y_test, y_pred))
print("\n")
print(classification_report(y_test, y_pred))
```

0.728

```
[[ 15  59]
 [   9 167]]
```

	precision	recall	f1-score	support
0	0.62	0.20	0.31	74
1	0.74	0.95	0.83	176
micro avg	0.73	0.73	0.73	250
macro avg	0.68	0.58	0.57	250
weighted avg	0.71	0.73	0.68	250

In [ ]:

In [ ]: